

IEEE International Symposium on Information Theory, Los Angeles (virtual) 2020

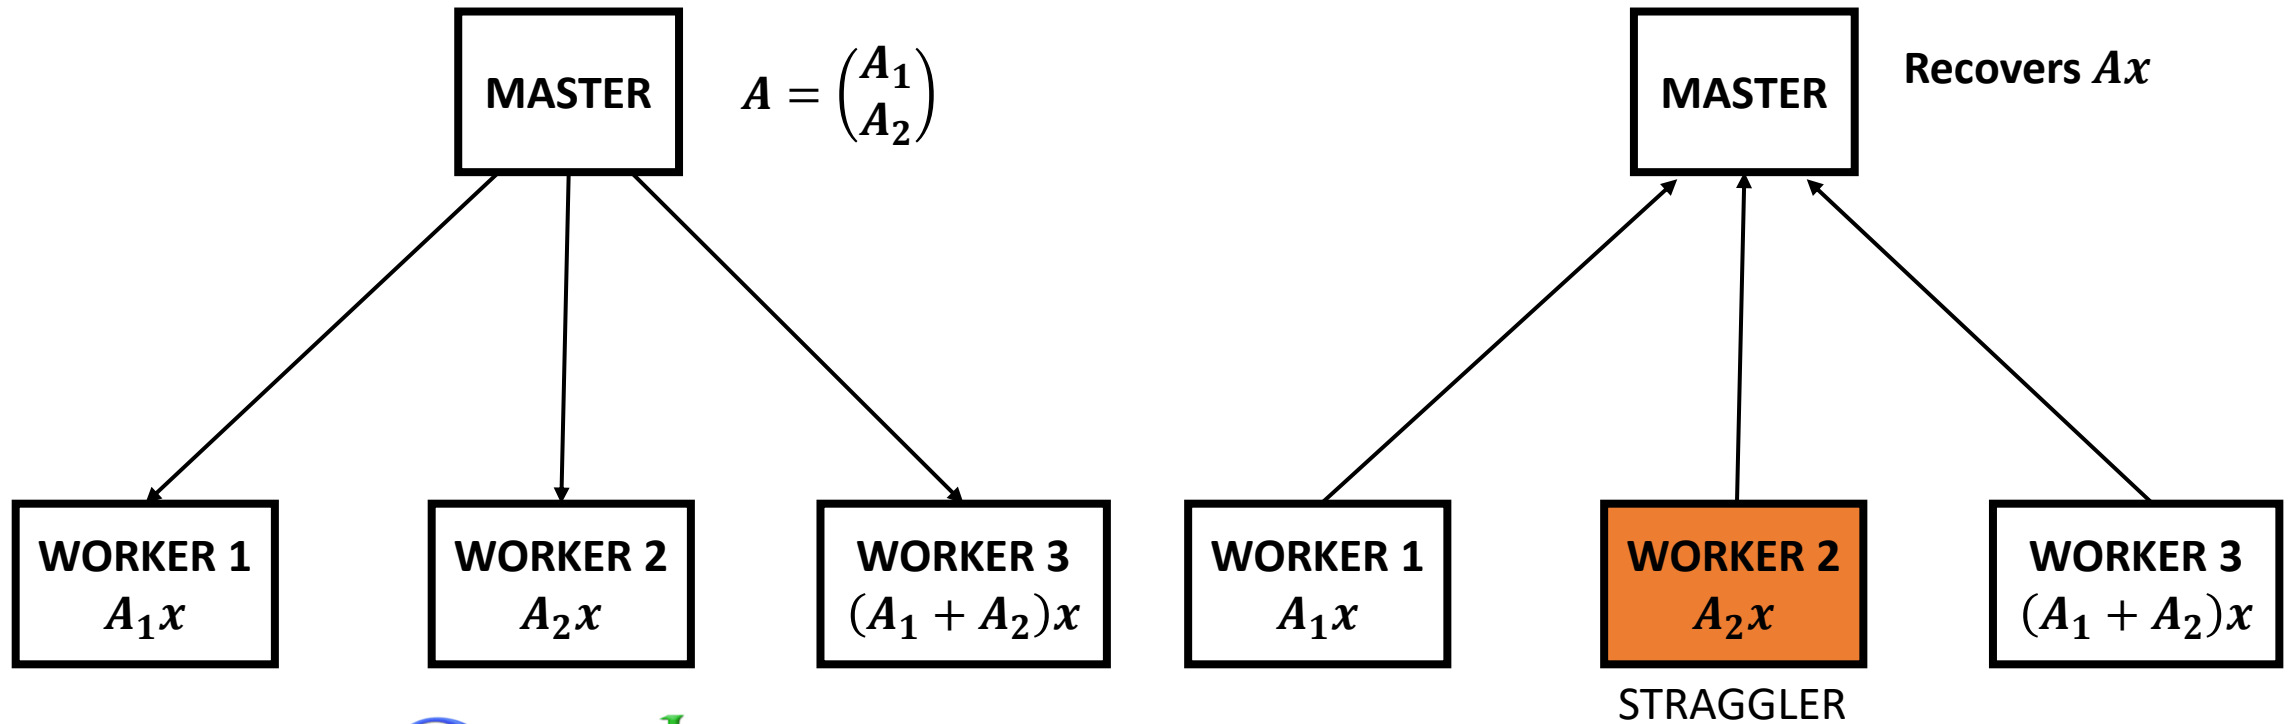
Optimizing the Transition Waste in Coded Elastic Computing

Hoang DAU

RMIT University (Melbourne)

Joint work with Ryan Gabrys, Yu-Chih Huang, Chen Feng, Quang-Hung Luu, Eidah Alzahrani, and Zahir Tari

Coded Computing Toy Example



Google


PageRank



The master can recover Ax with one straggler

Coded Elastic Computing

- $[L = 3, K = 2]$ coded computing scheme (CCS) in this example
- If there are more than three machines, must we change the CCS? **NO: Striping**



Coded Elastic Computing

	Machine 1	Machine 2	Machine 3	Machine 4	Machine 5
Task 1: A_1x	$A_{1,1}x$	$(A_{1,1} + A_{1,2})x$	$(A_{1,1} + 2A_{1,2})x$	$(A_{1,1} + 3A_{1,2})x$	$(A_{1,1} + 4A_{1,2})x$
Task 2: A_2x	$A_{2,1}x$	$(A_{2,1} + A_{2,2})x$	$(A_{2,1} + 2A_{2,2})x$	$(A_{2,1} + 3A_{2,2})x$	$(A_{2,1} + 4A_{2,2})x$
Task 3: A_3x	$A_{3,1}x$	$(A_{3,1} + A_{3,2})x$	$(A_{3,1} + 2A_{3,2})x$	$(A_{3,1} + 3A_{3,2})x$	$(A_{3,1} + 4A_{3,2})x$
Task 4: A_4x	$A_{4,1}x$	$(A_{4,1} + A_{4,2})x$	$(A_{4,1} + 2A_{4,2})x$	$(A_{4,1} + 3A_{4,2})x$	$(A_{4,1} + 4A_{4,2})x$
Task 5: A_5x	$A_{5,1}x$	$(A_{5,1} + A_{5,2})x$	$(A_{5,1} + 2A_{5,2})x$	$(A_{5,1} + 3A_{5,2})x$	$(A_{5,1} + 4A_{5,2})x$

- Machines store all “tasks”
- Machines only execute shaded tasks according to a Task Allocation Scheme (TAS)
- When a machine leaves the system, all tasks (completed/uncompleted) on that machine are lost
- When a machine leaves the system, the set of tasks allocated to existing machines must be updated

Coded Elastic Computing

- Why do we consider a variable number of machines?
 - Industries are offering **low-priority** virtual machines at a **reduced price** (up to 90% down): Amazon EC2 Spot, Microsoft Azure Batch
 - **Low priority**: e.g., a Spot instance can be pre-empted on short notice (2 minutes)
 - E.g., we rent 7 VMs = 3 on-demand + 4 Spot, to speed up our computation at low cost, need to make our **task assignment scheme** ready for **$N = 3,4,5,6,7$ machines**
- Referred to as coded “**elastic**” computing (Yang et al. ISIT’2019): elasticity means the number of machines varies
- Another follow-up work: extends the cyclic task assignment to machines with **varying computational speeds**

N. Woolsey, R-R. Chen, and M. Ji, “*Heterogeneous Computation Assignments in Coded Elastic Computing*”, 2020, <https://arxiv.org/abs/2001.04005>

Coded Elastic Computing

- In general, we divide \mathbf{A} into F equal-sized submatrices $\mathbf{A}_0, \dots, \mathbf{A}_{F-1}$
- The f -th *task/stripe* corresponds to the computation of $\mathbf{A}_f \mathbf{x}$ ($0 \leq f \leq F - 1$)
- Any $[L, K]$ CCS can be used for each task f
- To assign tasks/stripes to machines, we need a **Task Assignment Scheme** (TAS)

Definition: An $[N, L, F]$ -TAS is a list of sets $\mathcal{S}^N = (S_1^N, S_2^N, \dots, S_N^N)$, where

- $N = \#$ machines
- $S_n^N \subseteq [[F]] := \{0, 1, \dots, F - 1\}$ is the set of task indices assigned to Machine n
- **L -redundancy:** each index in $[[F]]$ lies in L sets in \mathcal{S}^N (each task is computed by L machines)
- **Load Balancing:** $|S_n^N| = \frac{LF}{N}$ (each machine computes $\frac{LF}{N}$ tasks)

Coded Elastic Computing

- $F = 20$ tasks (we have A_1, A_2, \dots, A_{20}), $N = 5$ machines, $L = 3$

S_1^5	S_2^5	S_3^5	S_4^5	S_5^5
$0 \rightarrow 3$			$0 \rightarrow 3$	$0 \rightarrow 3$
$4 \rightarrow 7$	$4 \rightarrow 7$			$4 \rightarrow 7$
$8 \rightarrow 11$	$8 \rightarrow 11$	$8 \rightarrow 11$		
	$12 \rightarrow 15$	$12 \rightarrow 15$	$12 \rightarrow 15$	
		$16 \rightarrow 19$	$16 \rightarrow 19$	$16 \rightarrow 19$

Cyclic Task Assignment when there are $N = 5$ machines (Yang *et al.* ISIT'19)

Coded Elastic Computing

- $F = 20$ tasks, $N = 4$ machines, $L = 3$

S_1^4	S_2^4	S_3^4	S_4^4
$0 \rightarrow 4$		$0 \rightarrow 4$	$0 \rightarrow 4$
$5 \rightarrow 9$	$5 \rightarrow 9$		$5 \rightarrow 9$
$10 \rightarrow 14$	$10 \rightarrow 14$	$10 \rightarrow 14$	
	$15 \rightarrow 19$	$15 \rightarrow 19$	$15 \rightarrow 19$

Cyclic Task Assignment when there are $N = 4$ machines (Yang *et al.* ISIT'19)

- When $N = 3 = L$, each machine computes **ALL** tasks

Transition Waste in Coded Elastic Computing

- **Transition**: when machines are **removed** or **added**
- **Transition waste**: Re-assigning tasks causes waste in computation

Definition

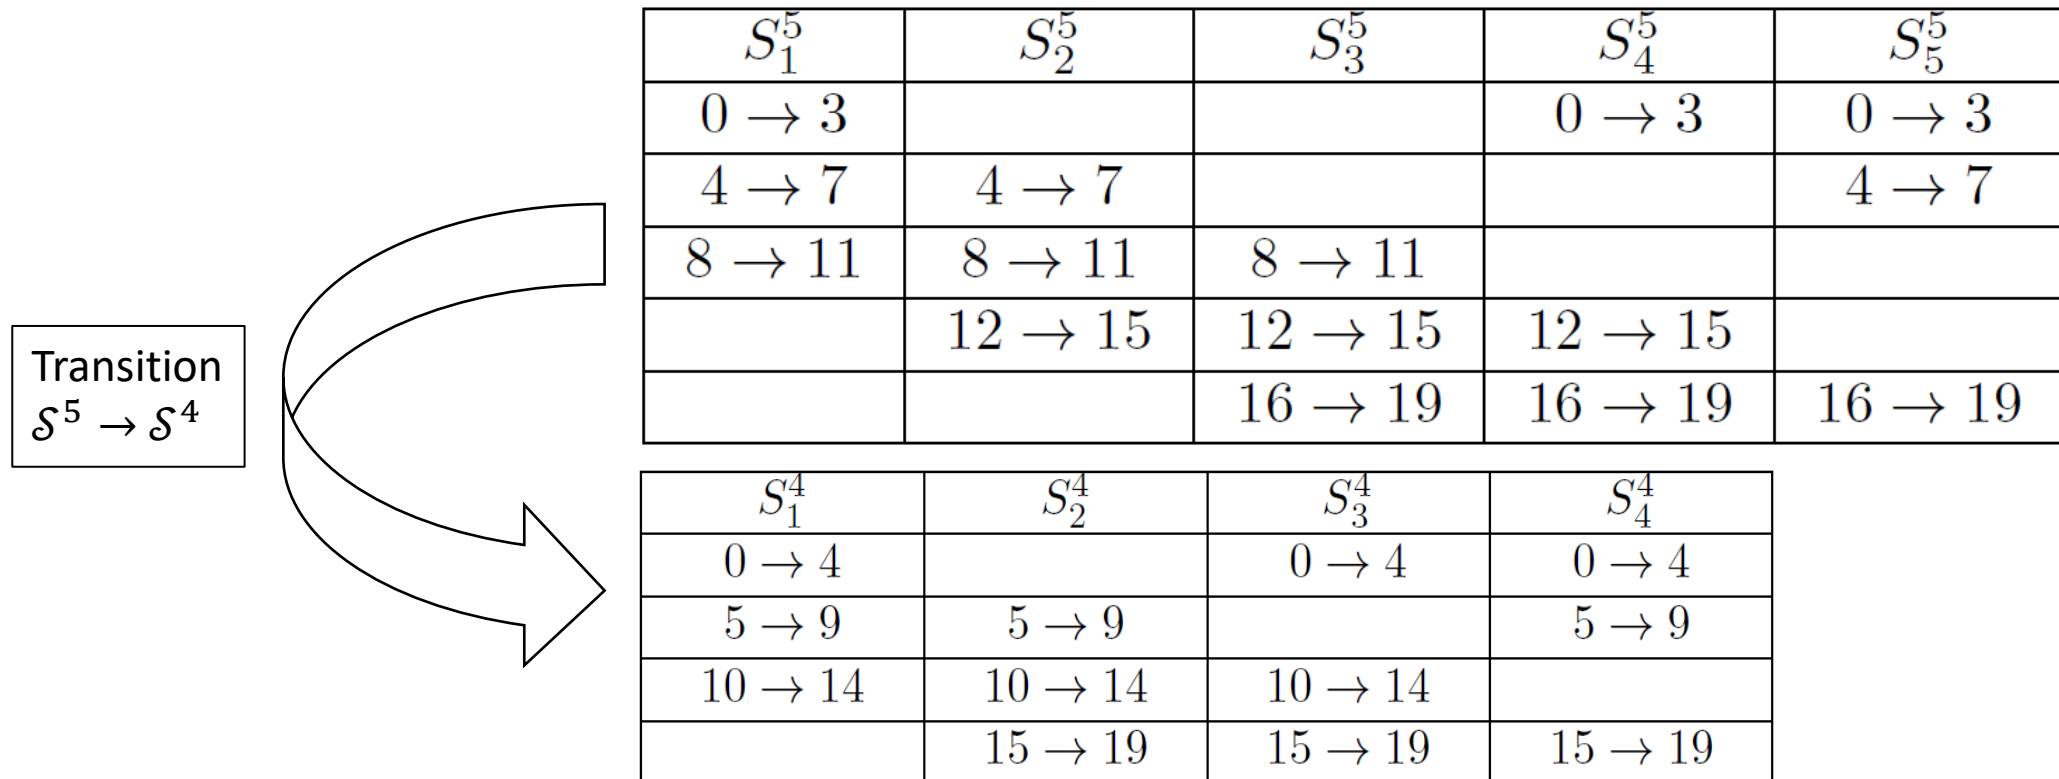
- TW at Machine n : $W_n = \text{\#abandoned tasks} + \text{\#new tasks} - \text{\#necessary change}$

Example

- Machine 3 was working on tasks $\{8 \rightarrow 19\}$,
- After the removal of Machine $n^* = 5$, it now works on tasks $\{0 \rightarrow 4\} \cup \{10 \rightarrow 19\}$
- It must **abandon two tasks** $\{8, 9\}$ and **takes over five new tasks** $\{0 \rightarrow 4\}$ (**symmetric difference**)
- The transition waste at Machine 3: $W_3 = 2 + 5 - 3 = 4$ tasks (3 is the necessary increase in the #tasks each machine takes as Machine 5 leaves, so not counted as waste)

Transition Waste in Coded Elastic Computing

- TW of a transition $\mathcal{S}^N \rightarrow \mathcal{S}^{N-1}$ or $\mathcal{S}^N \rightarrow \mathcal{S}^{N+1}$ is the **sum of TW at each machine**
- $W(\mathcal{S}^5 \rightarrow \mathcal{S}^4) = (0 + 3 - 3) + (1 + 4 - 3) + (2 + 5 - 3) + (3 + 6 - 3) = 12$



Transition Waste in Coded Elastic Computing

Our 1st contributions

- **Explicit formulas for transition wastes** in the cyclic TAS proposed in Yang *et al.*'19

Theorem 1. *The transition waste when transitioning from a cyclic (N, L, F) -TAS $\mathcal{S}_{\text{cyc}}^N$ to a cyclic $(N+1, L, F)$ -TAS $\mathcal{S}_{\text{cyc}}^{N+1}$ (defined in (2)) is given below (assuming $N > L$).*

$$W(\mathcal{S}_{\text{cyc}}^N \rightarrow \mathcal{S}_{\text{cyc}}^{N+1}) = \frac{N-1}{N+1}F.$$

Theorem 2. *The transition waste when Machine $n^* \in [N]$ leaves and the system transitions from a cyclic (N, L, F) -TAS $\mathcal{S}_{\text{cyc}}^N$ to a cyclic $(N-1, L, F)$ -TAS $\mathcal{S}_{\text{cyc}}^{N-1}$ (defined in (2)) is given as follows (assuming $N > L+1$).*

If $n^ < N-L$, $W_{n^*}(\mathcal{S}_{\text{cyc}}^N \rightarrow \mathcal{S}_{\text{cyc}}^{N-1})$ is*

$$((n^*-1)(n^*-2) + (N-L-n^*)(N-L-n^*+1)) \frac{F}{N(N-1)}.$$

If $n^ \geq N-L$, $W_{n^*}(\mathcal{S}_{\text{cyc}}^N \rightarrow \mathcal{S}_{\text{cyc}}^{N-1})$ is*

$$(n^*-1)(n^*-2) \frac{F}{N(N-1)}.$$

Transition Waste in Coded Elastic Computing

Our 2nd contributions

- Propose “**shifted**” cyclic TAS, which reduces the transition waste significantly

Ordinary cyclic TAS Start from index 0	S_1^4	S_2^4	S_3^4	S_4^4
	$0 \rightarrow 4$		$0 \rightarrow 4$	$0 \rightarrow 4$
	$5 \rightarrow 9$	$5 \rightarrow 9$		$5 \rightarrow 9$
	$10 \rightarrow 14$	$10 \rightarrow 14$	$10 \rightarrow 14$	
		$15 \rightarrow 19$	$15 \rightarrow 19$	$15 \rightarrow 19$

Shifted cyclic TAS Start from index i	S_1^4	S_2^4	S_3^4	S_4^4
	$17 \rightarrow 1$		$17 \rightarrow 1$	$17 \rightarrow 1$
	$2 \rightarrow 6$	$2 \rightarrow 6$		$2 \rightarrow 6$
	$7 \rightarrow 11$	$7 \rightarrow 11$	$7 \rightarrow 11$	
		$12 \rightarrow 16$	$12 \rightarrow 16$	$12 \rightarrow 16$

Transition Waste in Coded Elastic Computing

Our 2nd contributions

- Propose “shifted” cyclic TAS, which reduces the transition waste significantly
- Identify the shifted cyclic TAS with a **minimal TW** (under a divisibility condition)

Theorem 3. The transition waste when transitioning from a δ' -shifted cyclic (N, L, F) -TAS $\mathcal{S}_{\delta'-\text{cyc}}^N$ to a δ -shifted cyclic $(N + 1, L, F)$ -TAS $\mathcal{S}_{\delta-\text{cyc}}^{N+1}$ with $\delta = \delta' + \lfloor \frac{N+L-1}{2} \rfloor \frac{F}{N(N+1)}$ is

$$W(\mathcal{S}_{\delta'-\text{cyc}}^N \rightarrow \mathcal{S}_{\delta-\text{cyc}}^{N+1}) = \begin{cases} \frac{(N-L-1)(N-L+1)F}{2N(N+1)}, & \text{for odd } N-L, \\ \frac{(N-L)^2 F}{2N(N+1)}, & \text{for even } N-L. \end{cases}$$

Theorem 4. The transition waste when transitioning from a δ' -shifted cyclic (N, L, F) -TAS $\mathcal{S}_{\delta'-\text{cyc}}^N$ to a δ -shifted cyclic $(N - 1, L, F)$ -TAS $\mathcal{S}_{\delta-\text{cyc}}^{N-1}$ with $\delta = \delta' + (N - n^*) - \lfloor \frac{(N+L-2)}{2} \rfloor \frac{F}{N(N-1)}$, where Machine n^* leaves, is

$$W(\mathcal{S}_{\delta'-\text{cyc}}^N \rightarrow \mathcal{S}_{\delta-\text{cyc}}^{N-1}) = \begin{cases} \frac{(N-L-1)^2 F}{2N(N-1)}, & \text{for odd } N-L, \\ \frac{(N-L)(N-L-2)F}{2N(N-1)}, & \text{for even } N-L. \end{cases}$$

Transition Waste in Coded Elastic Computing


Our 3rd contributions

- How about **zero-waste** transitions?
- Theorem 6: \exists a zero-waste transition $\Leftrightarrow \exists$ a perfect matching in transition graph
- Intuition from Theorem 6: **intersections of sets** in the TAS should be **small**
- This led to the idea of using a **configuration** from combinatorial designs

Definition: A (v, k) (symmetric) configuration is an incident structure of v points and v lines such that

- Each line contains k points
- Each point lies on k lines
- Two different points lie on at most one line

Example: The Fano plane is a $(7, 7)$ configuration.



Transition Waste in Coded Elastic Computing

Our 3rd contributions

- How about **zero-waste** transitions?

S_1^7	S_2^7	S_3^7	S_4^7	S_5^7	S_6^7	S_7^7
{0, 1}	{0, 1}	{0, 1}				
{2, 3}			{2, 3}	{2, 3}		
{4, 5}					{4, 5}	{4, 5}
	{6, 7}		{6, 7}			{6, 7}
	{8, 9}			{8, 9}	{8, 9}	
		{10, 11}	{10, 11}		{10, 11}	
		{12, 13}		{12, 13}		{12, 13}

Example: using Fano plane, we can construct a TAS with 7 machines, which allows zero-waste transitions within 5, 6, and 7 machines.

Transition Waste in Coded Elastic Computing

Our 3rd contributions

- **Theorem 7**: \exists an (N_{\max}, L) configuration $\implies \exists$ zero-waste transitions between N_{\min} and N_{\max} machines, where $N_{\min} \approx \frac{N_{\max}}{2}$ from standard constructions of configurations $(L \approx \sqrt{N_{\max}})$

Open problem: does there exist zero-waste transitions for **arbitrary** L, N_{\max}, N_{\min} ?

Open problem: transition waste taking into account the sets of completed tasks?

Open problem: min-max transition waste?